

ADVANCED BAYESIAN MODELING - ASSIGNMENT 3

Shyam Shah
14/03/2021

1 (a)

The code to simulate the posterior for μ and σ^2 and produce an autocorrelation plot is below.

```
## Gibbs sampler for (partially conjugate) analysis of Flint data

n <- 271
ybar <- 1.48
s.2 <- 1.684

mu0 <- 1.10
tau.2.0 <- sigma.2.0 <- 1.17
nu0 <- 1

mun <- function(sigma.2){
  (mu0/tau.2.0 + n*ybar/sigma.2) / (1/tau.2.0 + n/sigma.2)
}

tau.2.n <- function(sigma.2){
  1 / (1/tau.2.0 + n/sigma.2)
}

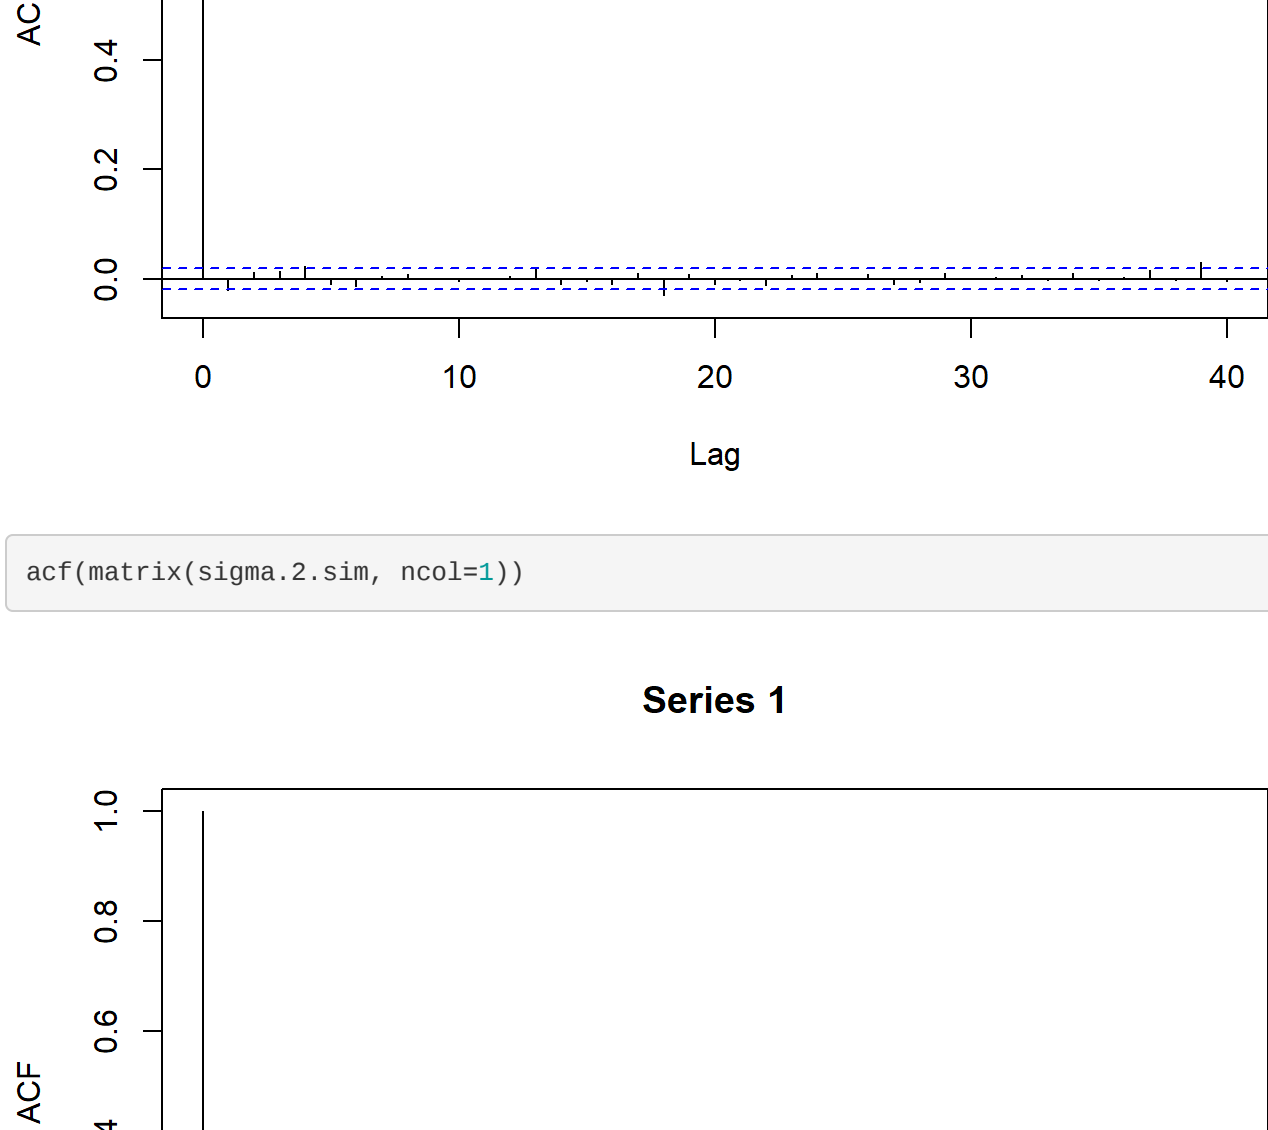
sigma.2.n <- function(mu){
  (nu0*sigma.2.0 + (n-1)*s.2 + n*(mu-ybar)^2) / (nu0 + n)
}

n.sim <- 10000
mu.sim <- numeric(n.sim)
sigma.2.sim <- numeric(n.sim)

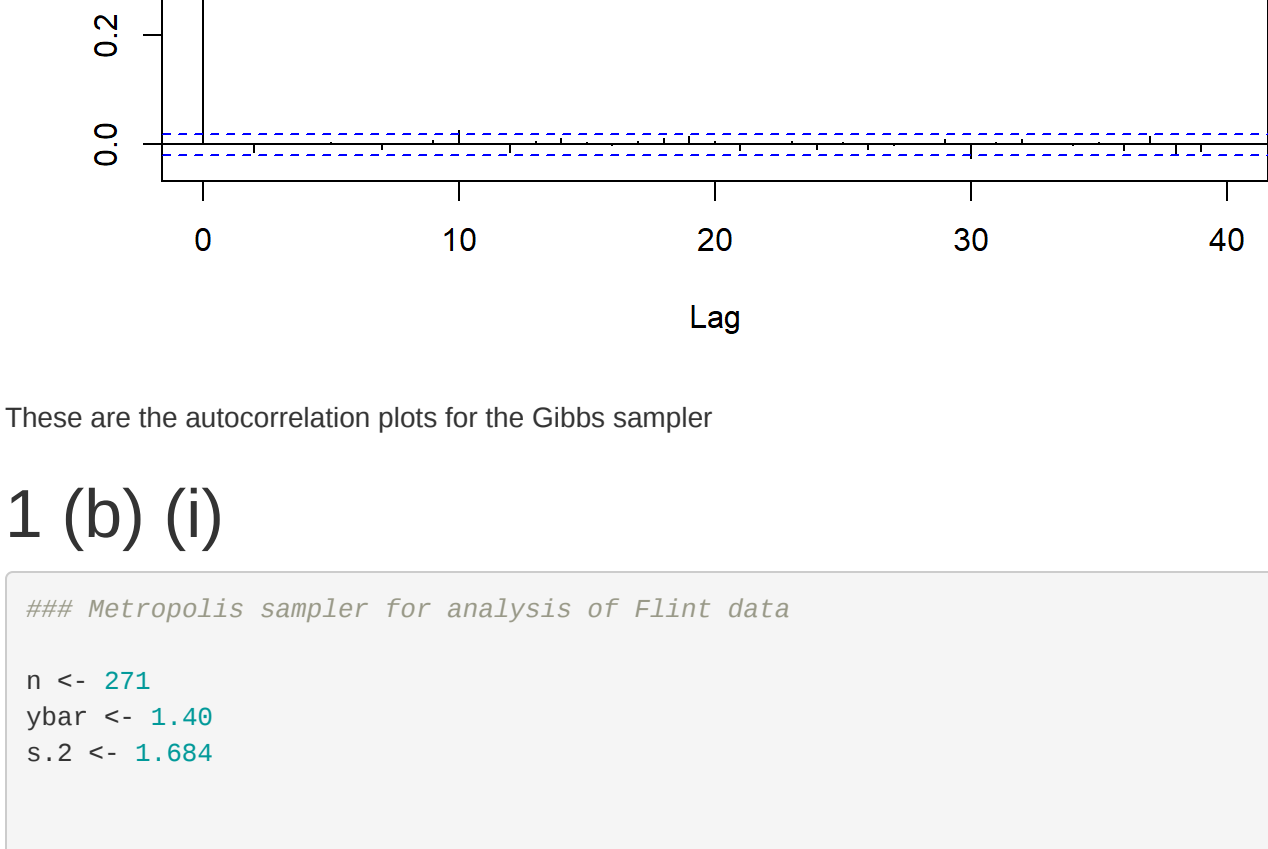
mu.sim[1] <- 1.4 # starting value
sigma.2.sim[1] <- 1.7 # starting value

for(t in 2:n.sim){
  mu.sim[t] <- rnorm(1, mun(sigma.2.sim[t-1]),
    sqrt(tau.2.n(sigma.2.sim[t-1])))
  sigma.2.sim[t] <- 1 / rgamma(1, (nu0+n)/2,
    (nu0+n)*sigma.2.n(mu.sim[t])/2)
}

acf(matrix(mu.sim, ncol=1))
```



```
acf(matrix(sigma.2.sim, ncol=1))
```



These are the autocorrelation plots for the Gibbs sampler

1 (b) (i)

```
## Metropolis sampler for analysis of Flint data

n <- 271
ybar <- 1.48
s.2 <- 1.684

mu0 <- 1.10
tau.2.0 <- sigma.2.0 <- 1.17
nu0 <- 1

dinvchisq <- function(x, df, scalesq){ # inverse chi-square density
  if(x>0){
    ((df/2)^(df/2) / gamma(df/2)) * sqrt(scalesq)*df * x^(-(df/2+1)) *
    exp(-df*scalesq/(2*x))
  } else 0
}

likelihood <- function(mu, sigma.2){
  sigma.2^(n/2) * exp(-(n-1)*s.2/(2*sigma.2)) *
  exp(-n*(mu-ybar)^2/(2*sigma.2))
}

ratio <- function(mu.prop, sigma.2.prop, mu.old, sigma.2.old){
  dnorm(mu.prop, mu0, sqrt(tau.2.0)) * dinvchisq(sigma.2.prop, nu0, sigma.2.0) *
  likelihood(mu.prop, sigma.2.prop) /
  (dnorm(mu.old, mu0, sqrt(tau.2.0)) * dinvchisq(sigma.2.old, nu0, sigma.2.0) *
  likelihood(mu.old, sigma.2.old))
}

n.sim <- 10000
mu.sim <- numeric(n.sim)
sigma.2.sim <- numeric(n.sim)
accept.prob <- numeric(n.sim-1)

rho <- 0.03

mu.sim[1] <- 1.4 # starting value
sigma.2.sim[1] <- 1.7 # starting value

for(t in 2:n.sim){
  mu.prop <- rnorm(1, mu.sim[t-1], sqrt(rho))
  sigma.2.prop <- rnorm(1, sigma.2.sim[t-1], sqrt(rho))

  accept.prob[t-1] <-
  min(ratio(mu.prop, sigma.2.prop, mu.sim[t-1], sigma.2.sim[t-1]), 1)
  if(runif(1) < accept.prob[t-1]){
    mu.sim[t] <- mu.prop
    sigma.2.sim[t] <- sigma.2.prop
  } else {
    mu.sim[t] <- mu.sim[t-1]
    sigma.2.sim[t] <- sigma.2.sim[t-1]
  }
}


mean(accept.prob)
```

```
## [1] 0.369908
```

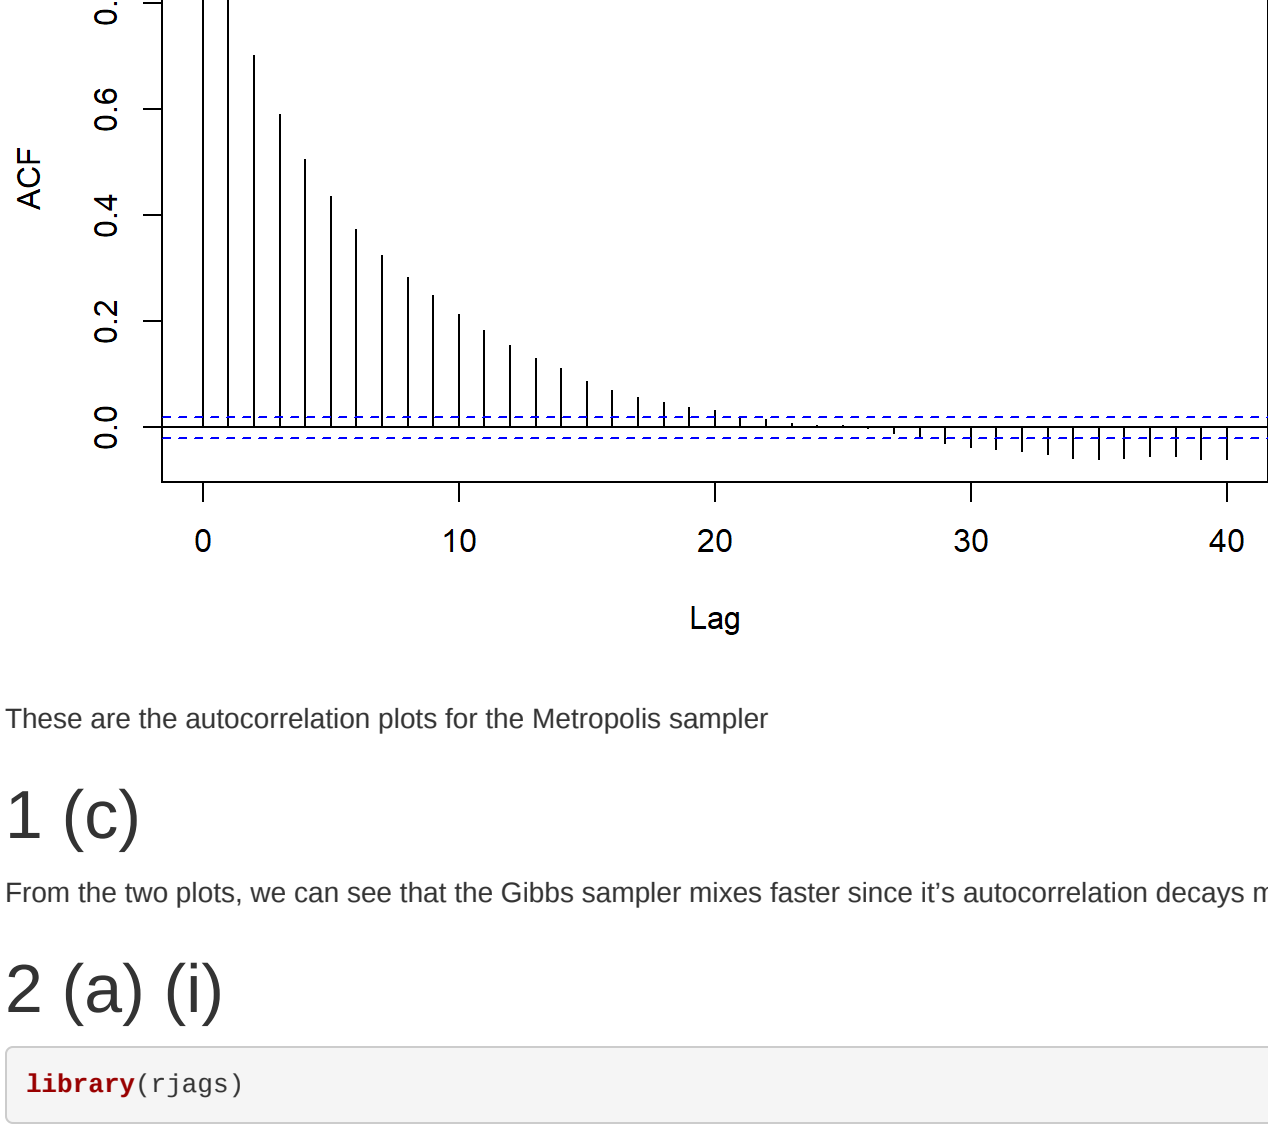
The value found for ρ is 0.03.

#1 (b) (ii)

```
acf(mu.sim)
```



```
acf(sigma.2.sim)
```



These are the autocorrelation plots for the Metropolis sampler

1 (c)

From the two plots, we can see that the Gibbs sampler mixes faster since its autocorrelation decays much faster.

2 (a) (i)

```
library(rjags)
```

```
## Loading required package: coda
```

```
## Linked to JAGS 4.3.0
```

```
## Loaded modules: basemod, bugs
```

```
# prepare the data
d <- read.table("./polls2016.txt", header=TRUE)
d$sigma <- d$ME/2

# create initialization lists
initial.vals1 <- list(list(mu=100, tau=0.01),
  list(mu=100, tau=0.01),
  list(mu=100, tau=100),
  list(mu=100, tau=100))

# create the model
m1 <- jags.model("./polls2016.bug", d, initial.vals1, n.chains=4)
```

```
## Warning in jags.model("./polls2016.bug", d, initial.vals1, n.chains = 4): Unused
## variable "ME" in data
```

```
## Warning in jags.model("./polls2016.bug", d, initial.vals1, n.chains = 4): Unused
## variable "ME" in data
```

```
## Compiling model graph
## Resolving undeclared variables
## Allocating nodes
## Graph information:
## Observed stochastic nodes: 7
## Unobserved stochastic nodes: 9
## Total graph size: 42
## Initializing model
```

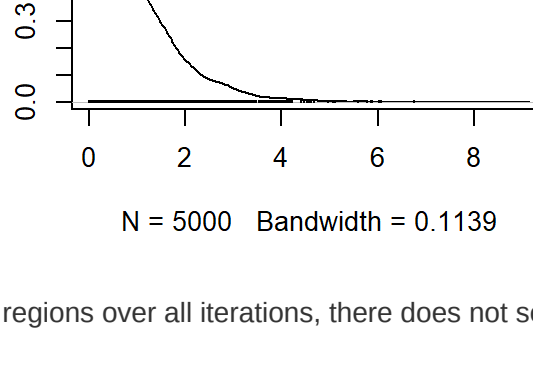
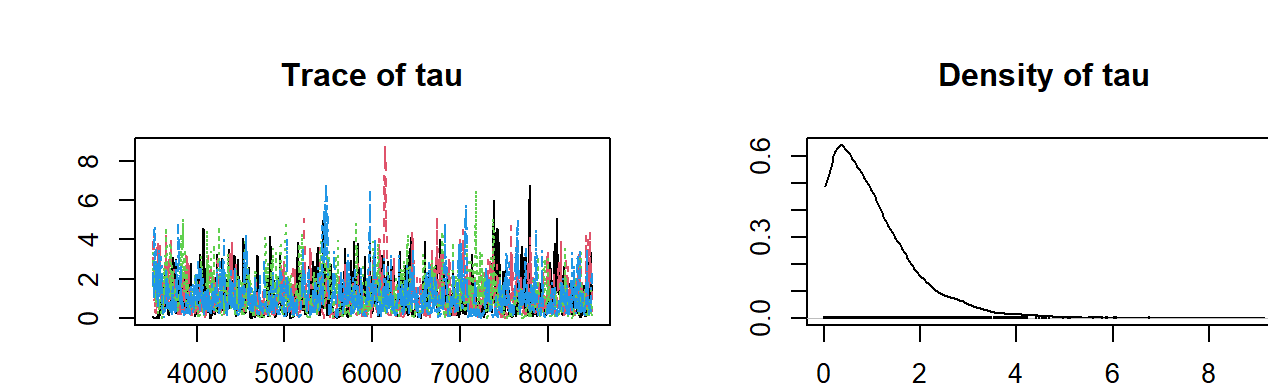

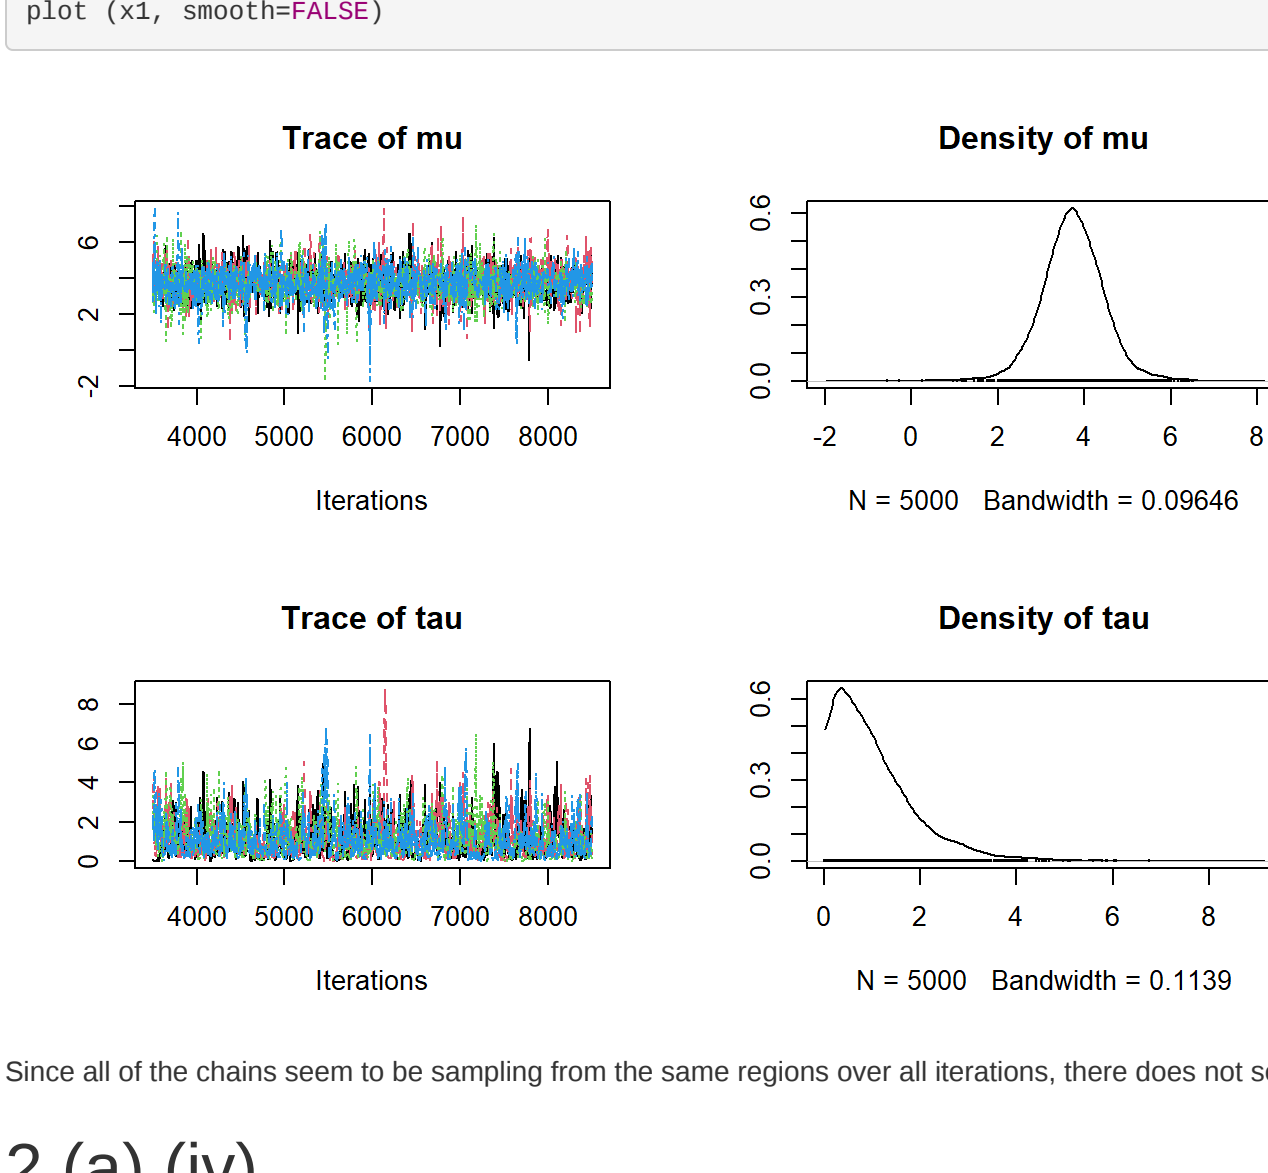
2 (a) (ii)

```
# 2500 iterations of burn in
update (m1, 2500)

# 5000 iterations of monitoring
x1 <- coda.samples(m1, c("mu", "tau"), n.iter=5000)
```

#2 (a) (iii)

```
# display trace plots
plot (x1, smooth=FALSE)
```



Since all of the chains seem to be sampling from the same regions over all iterations, there does not seem to be any convergence problems.

2 (a) (iv)


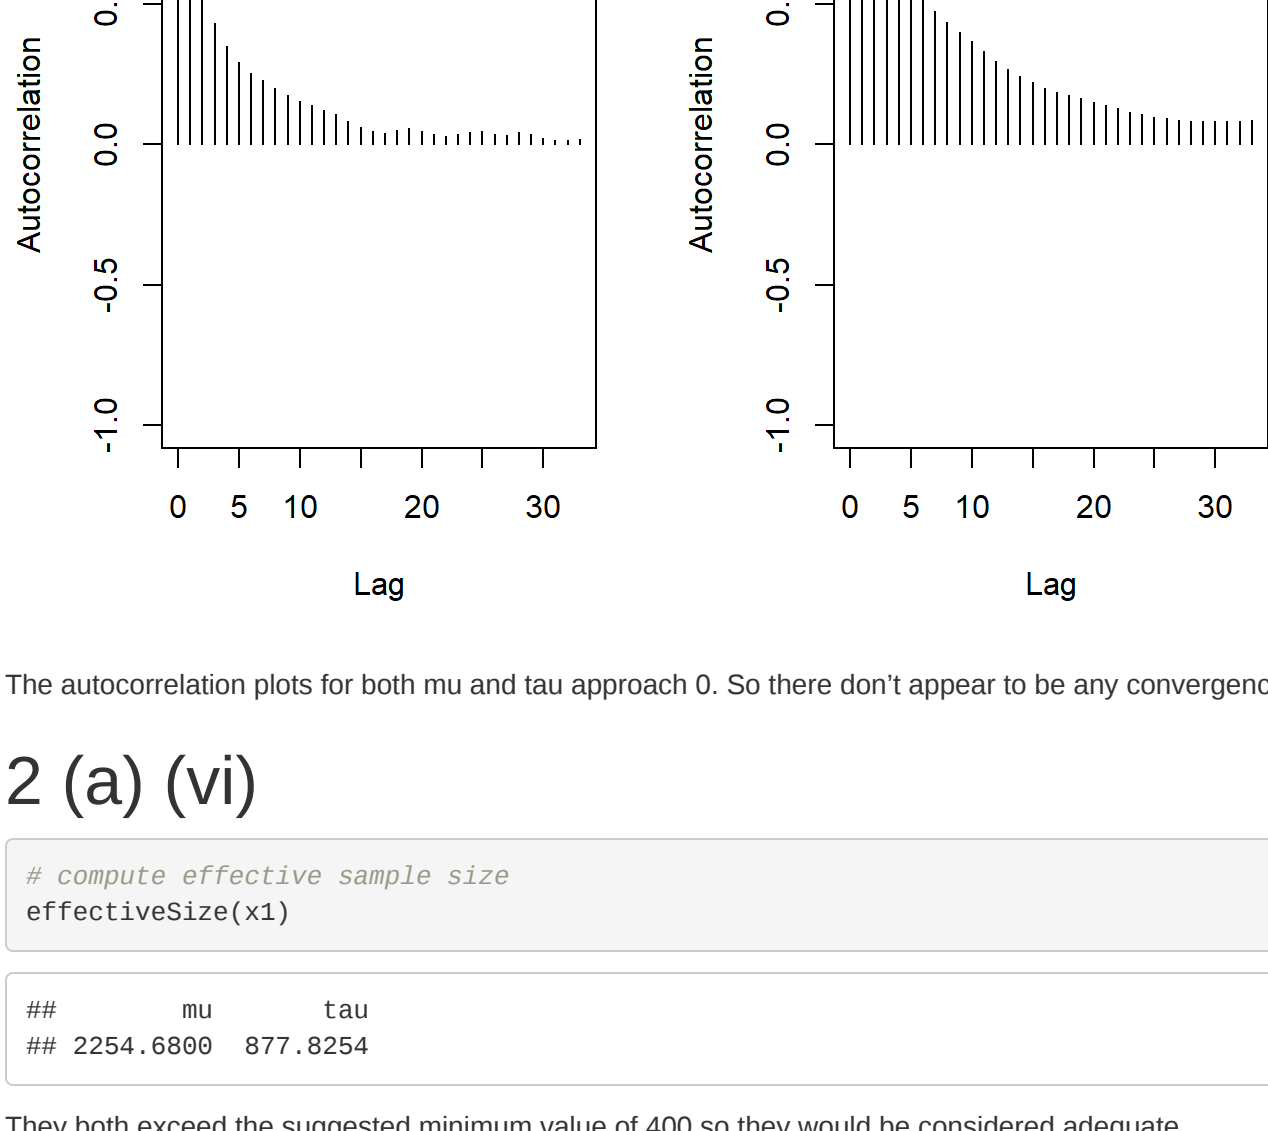
```
# get the Gelman-Rubin statistic
gelman.diag(x1, autoburnin=FALSE)
```

```
## Potential scale reduction factors:
##
## Point est. Upper C.I.
## mu 1.01 1.01
## tau 1.00 1.01
##
## Multivariate psrf
##
## 1.01
```

The values under point est. represent the Gelman-Rubin statistic and since they are very close to 1 (and both less than 1.1), this shows that there were no convergence problems.

2 (a) (v)

```
# create autocorrelation plot for first chain
autocorr.plot (x1[[1]])
```



The autocorrelation plots for both mu and tau approach 0. So there don't appear to be any convergence problems.

2 (a) (vi)

```
# compute effective sample size
effectiveSize(x1)
```

```
## mu tau
## 2254.6800 877.8254
```

They both exceed the suggested minimum value of 400 so they would be considered adequate.

2 (b) (i)

This is what the new model looks like:

```
model {
  for (j in 1:length(y)) {
    y[j] ~ dnorm(theta[j], 1/sigma[j]^2)
    theta[j] ~ dnorm(mu, 1/tau^2)
  }

  mu ~ dunif(-1000, 1000)
  tau <- exp(log(tau))
  logtau ~ dunif(-100, 100)
}
```

2 (b) (ii)

```
# create initialization lists with seeds because I saw lots of variator
initial.vals2 <- list(list(mu=100, logtau=log(0.01), .RNG.name="base:Wichmann-Hill1", .RNG.seed=1),
  list(mu=100, logtau=log(0.01), .RNG.name="base:Wichmann-Hill1", .RNG.seed=2),
  list(mu=100, logtau=log(100), .RNG.name="base:Wichmann-Hill1", .RNG.seed=3),
  list(mu=100, logtau=log(100), .RNG.name="base:Wichmann-Hill1", .RNG.seed=4))

# create the model
m2 <- jags.model("./polls20162.bug", d, initial.vals2, n.chains=4)
```

```
## Warning in jags.model("./polls20162.bug", d, initial.vals2, n.chains = 4):
## Unused variable "poll" in data
```

```
## Warning in jags.model("./polls20162.bug", d, initial.vals2, n.chains = 4):
## Unused variable "ME" in data
```

```
## Compiling model graph
## Resolving undeclared variables
## Allocating nodes
## Graph information:
## Observed stochastic nodes: 7
## Unobserved stochastic nodes: 9
## Total graph size: 44
## Initializing model
```

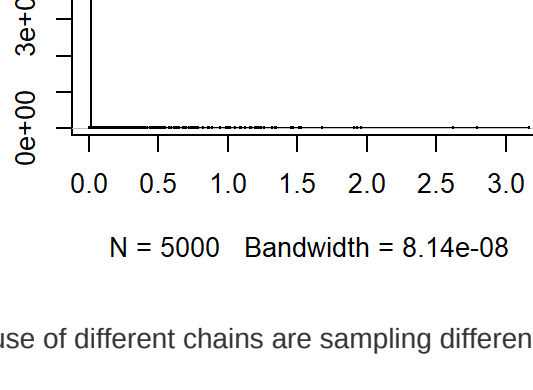

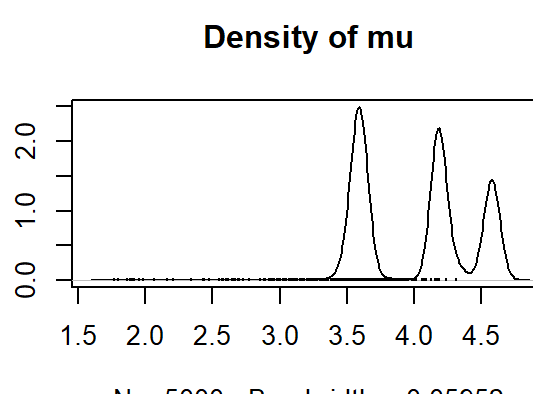
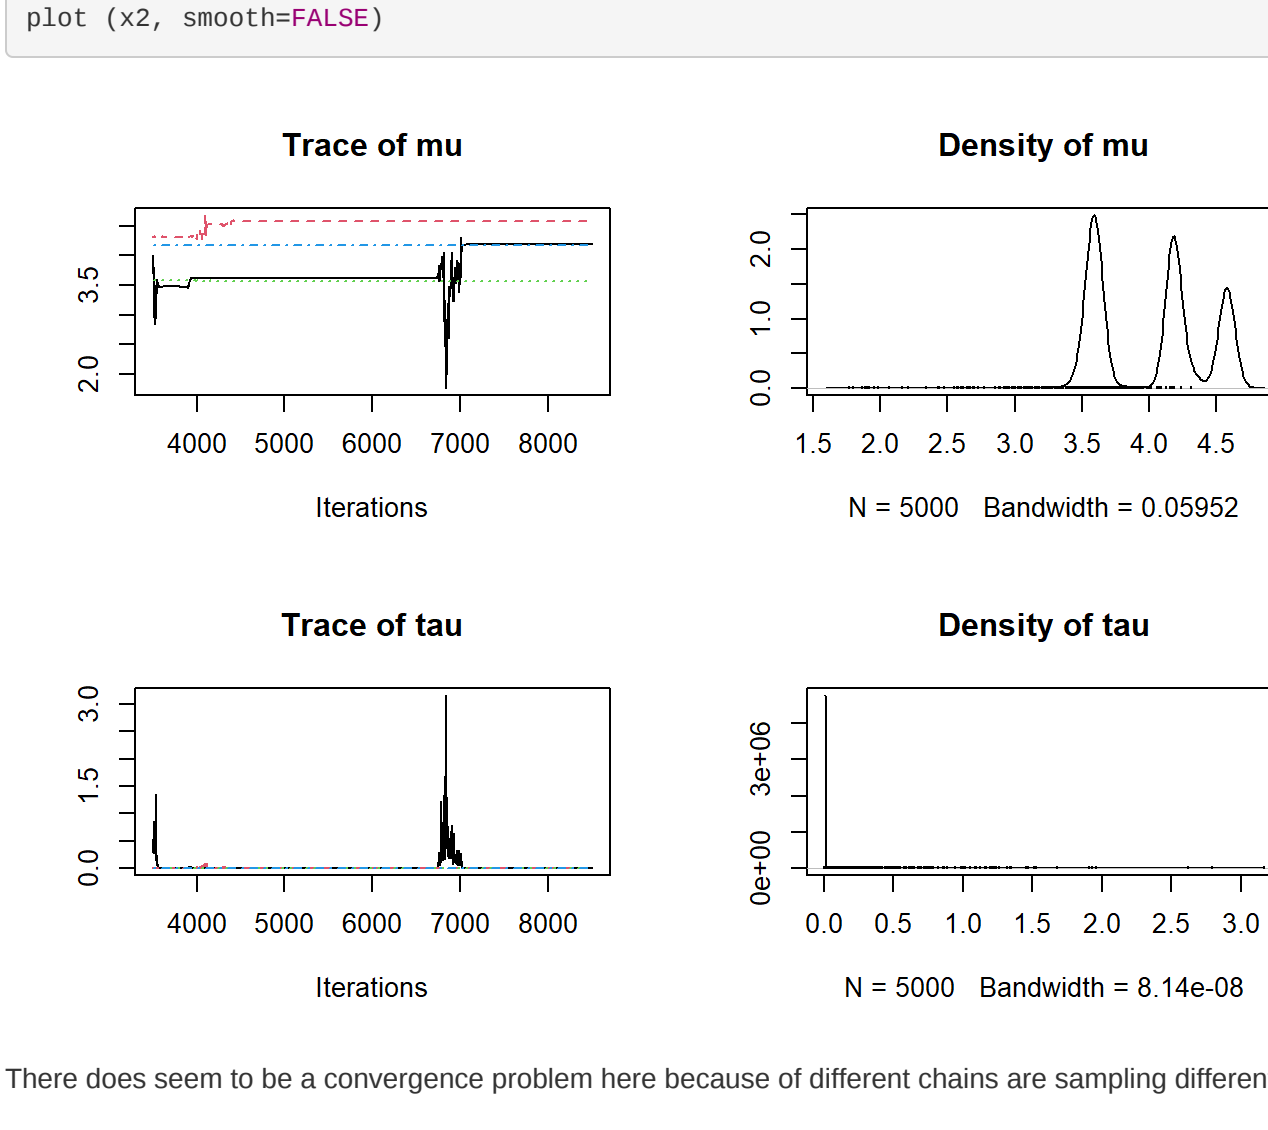
2 (b) (iii)

```
# 2500 iterations of burn in
update (m2, 2500)

# 5000 iterations of monitoring
x2 <- coda.samples(m2, c("mu", "tau"), n.iter=5000)
```

2 (b) (iv)

```
# display trace plots
plot (x2, smooth=FALSE)
```



There does seem to be a convergence problem here because of different chains are sampling different regions.

2 (b) (v)


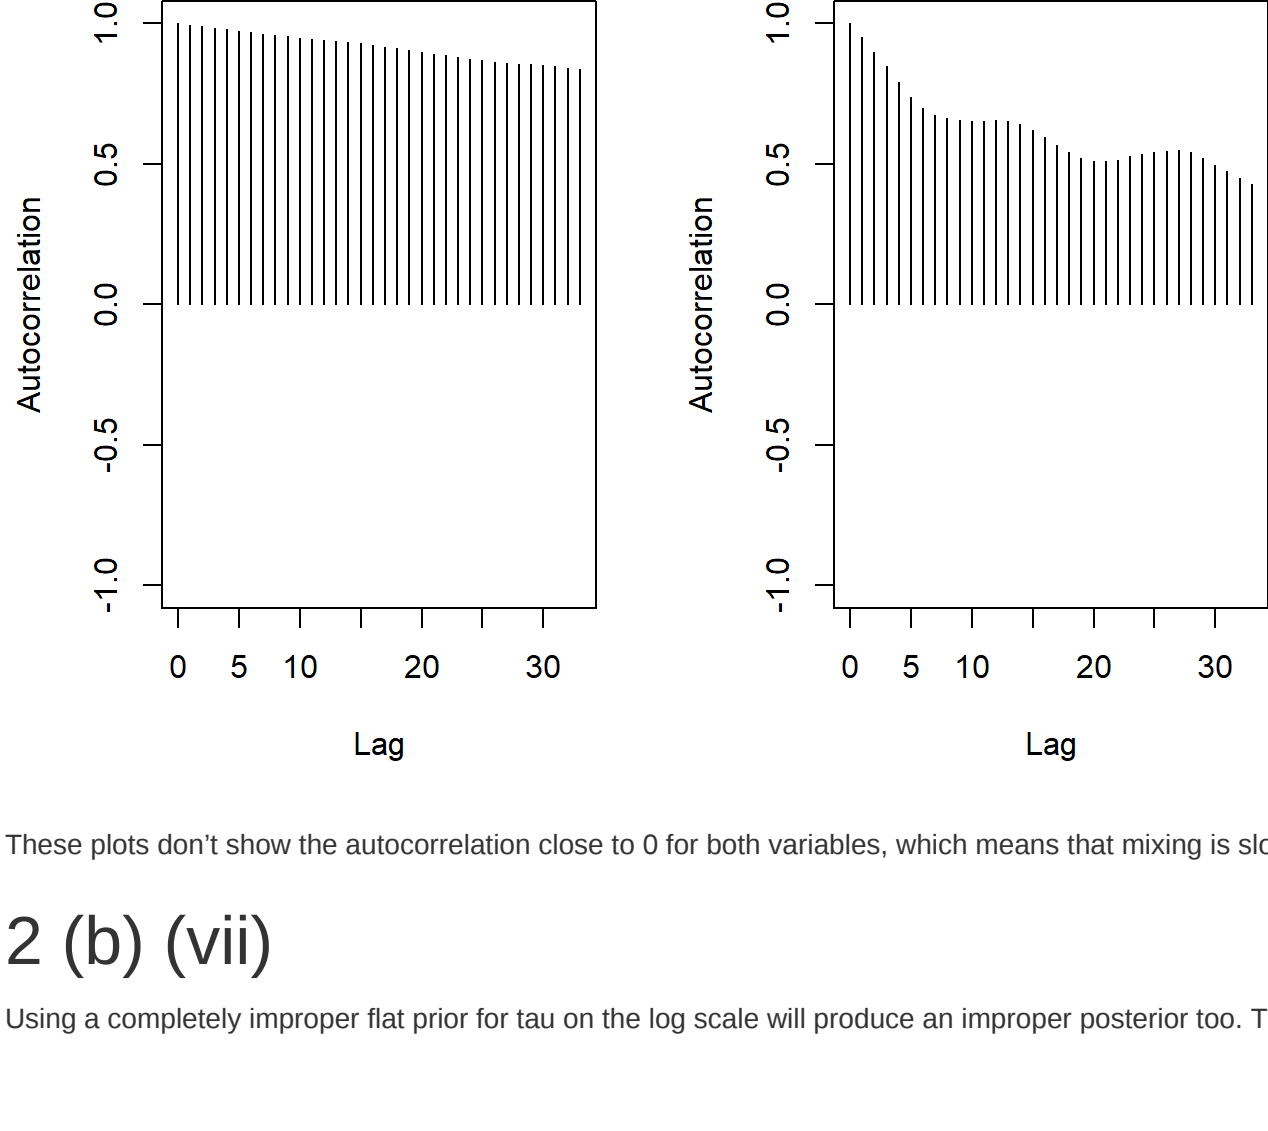
```
# get the Gelman-Rubin statistic
gelman.diag(x2, autoburnin=FALSE)
```

```
## Potential scale reduction factors:
##
## Point est. Upper C.I.
## mu 3.88 17.87
## tau 1.31 2.05
##
## Multivariate psrf
##
## 3.74
```

This also indicates an issue with convergence since the Gelman-Rubin statistic is over the suggested 1.1 for both variables.

2 (b) (vi)

```
# create autocorrelation plot for first chain
autocorr.plot (x2[[1]])
```



These plots don't show the autocorrelation close to 0 for both variables, which means that mixing is slow.

2 (b) (vii)

Using a completely improper flat prior for tau on the log scale will produce an improper posterior too. This is why we don't see convergence.