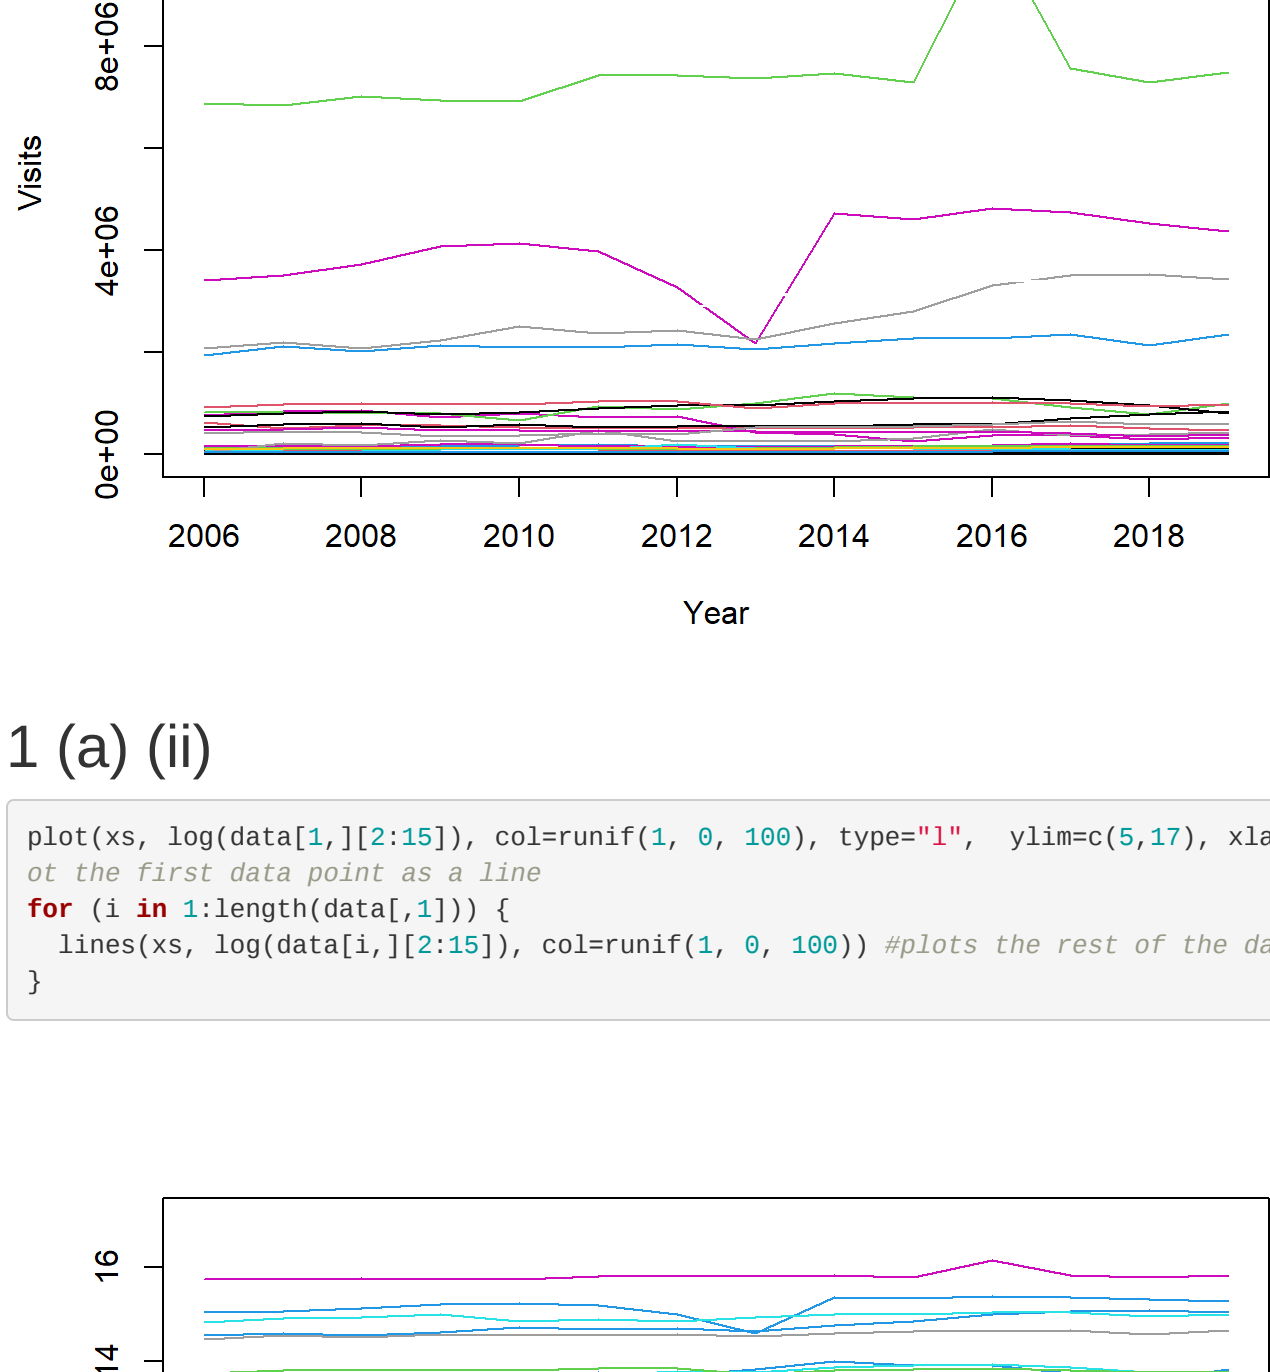


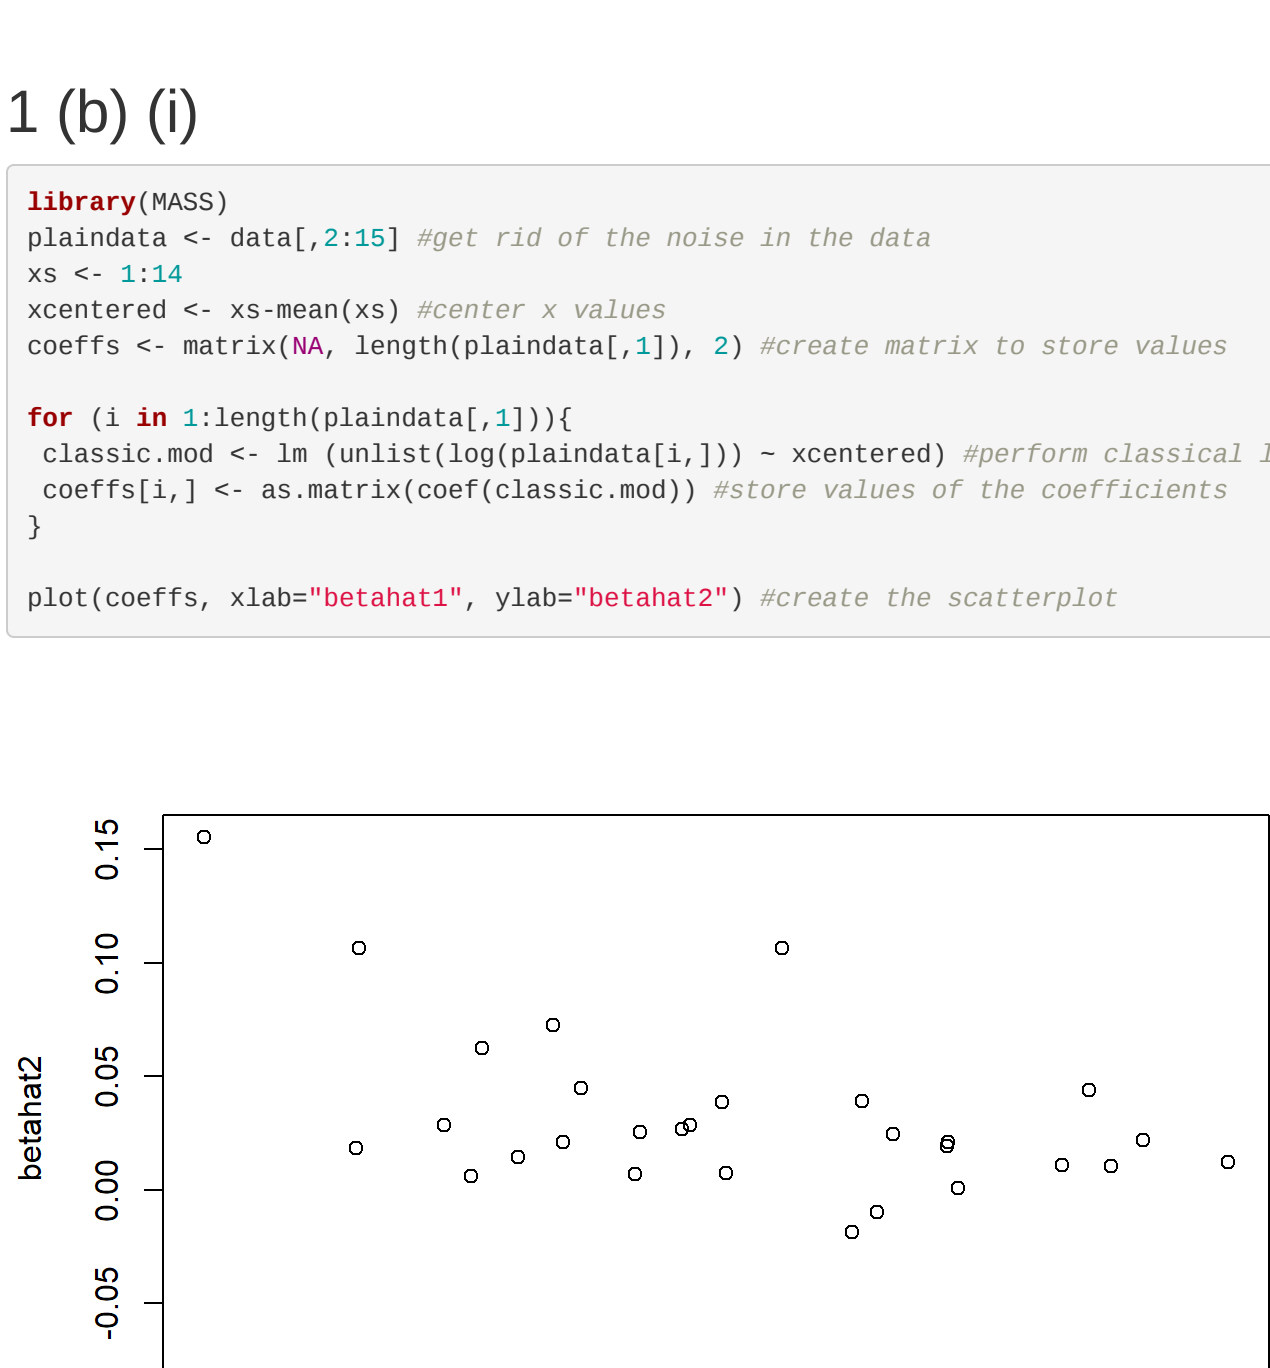
1 (a) (i)

```
data <- read.csv("../usparkkvisits.csv", header=TRUE) #read data
xs <- 2006:2019 #get list of years
plot(xs, data[,1][2:15], col=runif(1, 0, 100), type="l", ylim=c(0,1000000), xlab="Year", ylab="Visits") #plot t
# the first data point as a line
for (i in 1:length(data[,1])) {
  lines(xs, data[,1][2:15], col=runif(1, 0, 100)) #plots the rest of the data on the same graph
}
```



1 (a) (ii)

```
plot(xs, log(data[,1][2:15]), col=runif(1, 0, 100), type="l", ylim=c(5,17), xlab="Year", ylab="log(Visits)") #pl
# the first data point as a line
for (i in 1:length(data[,1])) {
  lines(xs, log(data[,1][2:15]), col=runif(1, 0, 100)) #plots the rest of the data on the same graph
}
```



1 (b) (i)

```
library(MASS)
plaindata <- data[,2:15] #get rid of the noise in the data
xs <- 1:14
xcentered <- xs-mean(xs) #center x values
coeffs <- matrix(NA, length(plaindata[,1]), 2) #create matrix to store values
for (i in 1:length(plaindata[,1])){
  classic.mod <- lm (unlist(log(plaindata[,1])) ~ xcentered) #perform classical linear regression
  coeffs[i,1] <- as.matrix(coeff(classic.mod)) #store values of the coefficients
}
```

```
plot(coeffs, xlab="betahat1", ylab="betahat2") #create the scatterplot
```



1 (b) (ii)

```
betahat2.mean <- mean(coeffs[,1])
betahat2.mean <- mean(coeffs[,2])
betahat2.mean
```

```
## [1] 12.18651
betahat2.mean
## [1] 0.028471
```

The sample mean for $\hat{\beta}_1^{(i)}$ is 12.18651 and the sample mean for $\hat{\beta}_2^{(i)}$ is 0.028471.

1 (b) (iii)

```
betahat2.var <- var(coeffs[,1])
betahat2.var <- var(coeffs[,2])
betahat2.var
```

```
## [1] 3.773097
betahat2.var
## [1] 0.001861814
```

The sample variance for $\hat{\beta}_1^{(i)}$ is 3.773097 and the sample variance for $\hat{\beta}_2^{(i)}$ is 0.001861814.

1 (b) (iv)

```
betahat.cov <- cov(coeffs[,1],coeffs[,2])
betahat.cov
```

```
## [1] -0.83864999
betahat.cov
```

The sample correlation between $\hat{\beta}_1^{(i)}$ and $\hat{\beta}_2^{(i)}$ is -0.83864999.

1 (c) (i)

```
#prepare data for the model
d1 <- list (visits = log(plaindata),
           year = xs,
           mubeta0 = c(0,0),
           Signamubetainv = rbind(c(0.000001, 0),
                                   c(5, 0.000001)),
           Sigma0 = rbind(c(100, 0),
                          c(0, 0.1)))

#initialize chains
initst <- list(list(sigmasyinv=10, mubeta=c(1000,1000),
                   Signamubetainv = rbind(c(100,0),
                                           c(0,100))),
              .RNG.name="base::Wichmann-Hill", .RNG.seed=123456),
              list(sigmasyinv=0.001, mubeta=c(-1000,1000),
                   Signamubetainv = rbind(c(100,0),
                                           c(5,100))),
              .RNG.name="base::Wichmann-Hill", .RNG.seed=123455),
              list(sigmasyinv=0, mubeta=c(1000,-1000),
                   Signamubetainv = rbind(c(0.001,0),
                                           c(0,0.001))),
              .RNG.name="base::Wichmann-Hill", .RNG.seed=123454),
              list(sigmasyinv=0.001, mubeta=c(-1000,-1000),
                   Signamubetainv = rbind(c(0.001,0),
                                           c(0,0.001))),
              .RNG.name="base::Wichmann-Hill", .RNG.seed=123453))

library(rjags)

## Loading required package: coda

## Linked to JAGS 4.3.0

## Loaded modules: base64,bugs

#print the model out
cat (readLines("../Jc1.bug"), sep= '\n')

## data {
##   dimv <- dim(visits)
##   yearcent <- year - mean(year)
## }
## model {
##   for (j in 1:dimv[1]) {
##     for (i in 1:dimv[2]) {
##       visits[j,i] ~ dnorm (beta[1,j]+beta[2,j]*yearcent[i], sigmasyinv)
##     }
##     beta[1:2,j] ~ dnmorm (mubeta, Signamubetainv)
##   }
##   mubeta ~ dnmorm(mubeta0, Signamubetainv)
##   Signamubetainv ~ dwnish(2*Sigma0, 2)
##   sigmasyinv ~ dgamma(0.0001, 0.0001)
## }
## SigmaBeta <- inverse(Signamubetainv)
## rho <- Signamubeta[1,2]/sqrt(Signamubeta[1,1]*Signamubeta[2,2])
## sigmasyq <- 1/sigmasyinv

#Create model
m1 <- jags.model("../Jc1.bug", d1, initst1, n.chains=4, n.adapt=1000)

## Compiling data graph
## Resolving undeclared variables
## Allocating nodes
## Initializing
## Reading data back into data table
## Compiling model graph
## Resolving undeclared variables
## Allocating nodes
## Graph information:
## Observed stochastic nodes: 428
## Unobserved stochastic nodes: 33
## Total graph size: 1408
## Initializing model

update(m1, 23000) # burn-in

x1 <- coda.samples(m1, c("mubeta", "Signamubeta", "sigmasyq", "rho"), n.iter=4000) #get samples
gelman.diag(x1, autoburnin=FALSE, multivariate=FALSE) #shows convergence if < 1.1

## Potential scale reduction factors:
##
## Point est. Upper C.I.
## Signamubeta[1,1] 1 1 1
## Signamubeta[2,1] 1 1
## Signamubeta[1,2] 1 1
## Signamubeta[2,2] 1 1
## mubeta[1] 1 1
## mubeta[2] 1 1
## rho 1 1
## sigmasy 1 1

effectiveSize(x1) #check effective size > 4000

## Signamubeta[1,1] Signamubeta[2,1] Signamubeta[1,2] Signamubeta[2,2] mubeta[1]
## 2527.22 14506.34 14506.34 14506.34 18287.62
## mubeta[2] rho sigmasy
## 15428.95 14486.81 12833.41
```

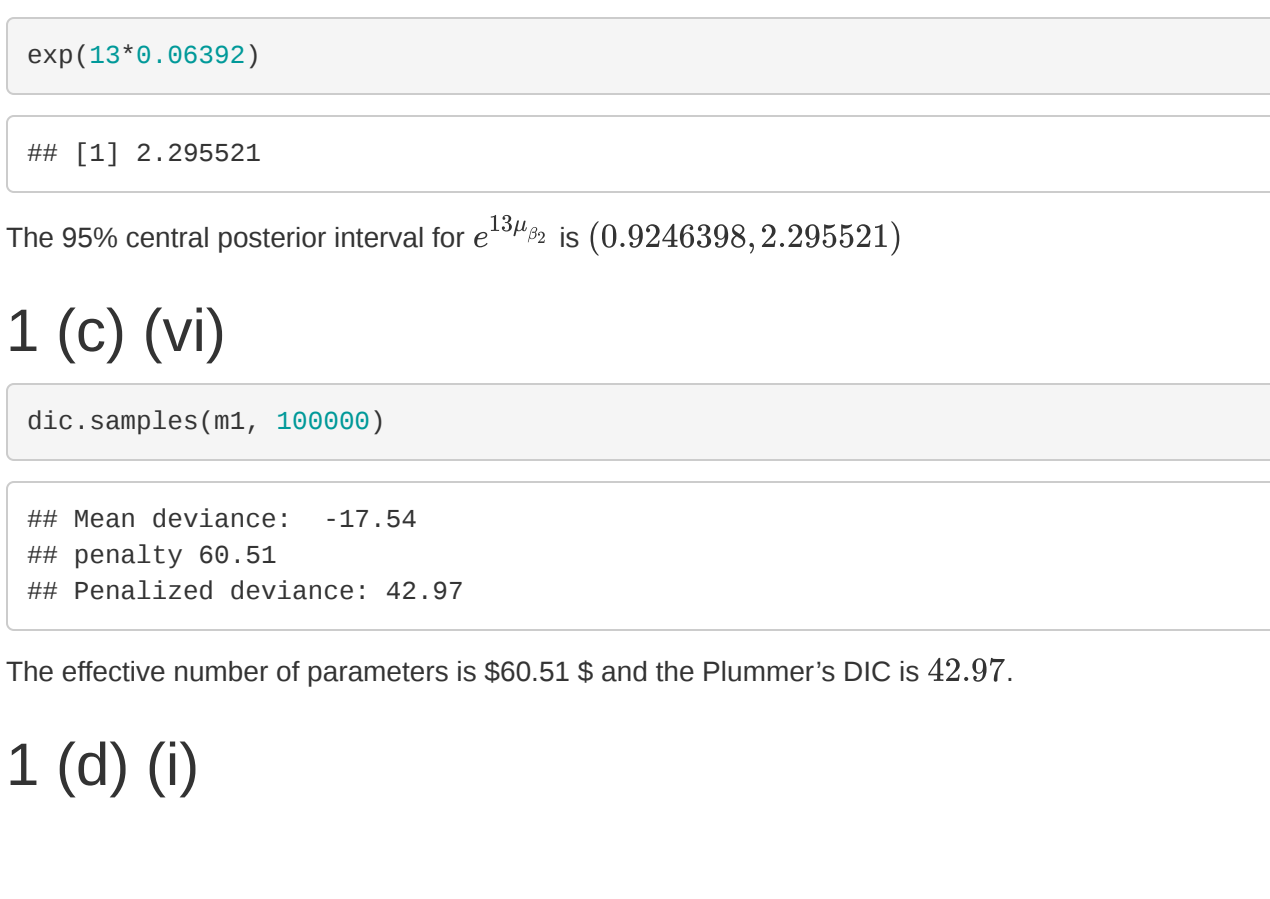
1 (c) (ii)

```
summary(x1)

##
## Iterations = 23081:27000
## Thinning interval = 1
## Number of chains = 4
## Sample size per chain = 4000
## 1. Empirical mean and standard deviation for each variable,
## plus standard error of the mean:
##
## Mean SD Naive SE Time-series SE
## Signamubeta[1,1] 11.050382 3.072853 2.429e-02 2.494e-02
## Signamubeta[2,1] -0.832060 0.802032 4.904e-04 5.106e-04
## Signamubeta[1,2] -0.832030 0.802032 4.904e-04 5.106e-04
## Signamubeta[2,2] 0.809202 0.802557 2.821e-05 2.886e-05
## mubeta[1] 12.129392 0.016129 4.823e-03 4.762e-03
## mubeta[2] 0.028392 0.017725 1.401e-04 1.402e-04
## rho -0.128954 0.177263 1.401e-03 1.475e-03
## sigmasy 0.056392 0.004254 3.353e-05 3.806e-05
##
## 2. Quantiles for each variable:
##
## 2.5% 25% 50% 75% 97.5%
## Signamubeta[1,1] 6.527671 8.873409 10.561764 1.262e+01 16.54533
## Signamubeta[2,1] -0.173227 -0.074900 -0.836568 -3.883e-06 0.07667
## Signamubeta[1,2] -0.173227 -0.074900 -0.836568 -3.883e-06 0.07667
## Signamubeta[2,2] 0.805458 0.807368 0.800776 1.059e-02 0.01527
## mubeta[1] -0.950000 11.769992 12.186512 1.200e+01 13.38096
## mubeta[2] -0.866027 0.016660 0.028313 3.997e-02 0.06392
## rho -0.455901 -0.242696 -0.123776 -1.981e-05 0.23164
## sigmasy 0.040780 0.053439 0.050101 5.912e-02 0.06538
```

1 (c) (iii)

```
densplot(x1[,c("rho")])
```



A 95% central posterior interval for ρ is $(-0.455901, 0.23164)$

1 (c) (iv)

```
mean(as.matrix(x1[,c("rho")]))<0 #posterior probability for rho < 0
## [1] 0.75
mean(as.matrix(x1[,c("rho")]))<0 & mean(as.matrix(x1[,c("rho")]))>0 #calculate Bayes factor as posterior/prior
## [1] 3
```

The posterior probability that $\rho < 0$ is 0.75, and the Bayes factor favoring $\rho < 0$ over $\rho > 0$ is 3 which means the data evidence is between barely noninformative and positive.

1 (c) (v)

```
# reading the quantiles for mubeta[2] from id11
exp(13*-0.006027)
## [1] 0.9246398
exp(13*-0.05393)
## [1] 2.295521
```

The 95% central posterior interval for $e^{13\rho}$ is $(0.9246398, 2.295521)$

1 (c) (vi)

```
dic.samples(m1, 100000)
## Mean deviance: -17.54
## penalty BIC: 51
## Penalized deviance: 42.97
```

The effective number of parameters is 560.51 \$ and the Plummer's DIC is 42.97.

1 (d) (i)

```
#prepare data for the model
d2 <- list (visits = log(plaindata),
           year = xs)

#initialize chains, adding initial values for each signamubeta and mubeta
initst2 <- list(list(sigmasyinv=10, mubeta1=1000, mubeta2=1000,
                   signamubeta1=1000, signamubeta2=1000,
                   .RNG.name="base::Wichmann-Hill", .RNG.seed=123456),
              list(sigmasyinv=0.001, mubeta1=-1000, mubeta2=-1000,
                   signamubeta1=1000, signamubeta2=1000,
                   .RNG.name="base::Wichmann-Hill", .RNG.seed=123457),
              list(sigmasyinv=10, mubeta1=1000, mubeta2=-1000,
                   signamubeta1=0.001, signamubeta2=0.001,
                   .RNG.name="base::Wichmann-Hill", .RNG.seed=123458),
              list(sigmasyinv=0.001, mubeta1=-1000, mubeta2=1000,
                   signamubeta1=0.001, signamubeta2=0.001,
                   .RNG.name="base::Wichmann-Hill", .RNG.seed=123459))

library(rjags)

#print the model out
cat (readLines("../Jd11.bug"), sep= '\n')

## data {
##   dimv <- dim(visits)
##   yearcent <- year - mean(year)
## }
## model {
##   for (j in 1:dimv[1]) {
##     for (i in 1:dimv[2]) {
##       visits[j,i] ~ dnorm (beta[1,j]+beta[2,j]*yearcent[i], sigmasyinv)
##     }
##     beta[1,j] ~ dnmorm (mubeta1, signamubeta1inv)
##     beta[2,j] ~ dnmorm (mubeta2, signamubeta2sqinv)
##   }
##   mubeta1 ~ dnmorm(0, 0.000001)
##   mubeta2 ~ dnmorm(0, 0.000001)
##   signamubeta1 ~ dunif(0, 1000)
##   signamubeta2 ~ dunif(0, 1000)
##   sigmasyinv ~ dgamma(0.0001, 0.0001)
## }
## signamubeta1sq <- signamubeta1^2
## signamubeta2sq <- signamubeta2^2
## signamubeta1sqinv <- 1/signamubeta1sq
## signamubeta2sqinv <- 1/signamubeta2sq
## sigmasyq <- 1/sigmasyinv

#Create model
m2 <- jags.model("../Jd11.bug", d2, initst2, n.chains=4, n.adapt=2000)

## Compiling data graph
## Resolving undeclared variables
## Allocating nodes
## Initializing
## Reading data back into data table
## Compiling model graph
## Resolving undeclared variables
## Allocating nodes
## Graph information:
## Observed stochastic nodes: 420
## Unobserved stochastic nodes: 65
## Total graph size: 1366
## Initializing model

update(m2, 80000) # burn-in

x2 <- coda.samples(m2, c("mubeta1", "mubeta2", "signamubeta1sq", "signamubeta2sq", "sigmasyq"), n.iter=4000) #get sam
ples
gelman.diag(x2, autoburnin=FALSE, multivariate=FALSE) #show convergence

## Potential scale reduction factors:
##
## Point est. Upper C.I.
## mubeta1 1.00 1.00
## mubeta2 1.00 1.00
## signamubeta1 1.00 1.27
## signamubeta2sq 1.00 1.00
## sigmasy 1.00 1.00

effectiveSize(x2) #check effective size > 4000

## mubeta1 mubeta2 signamubeta1sq signamubeta2sq sigmasy
## 15375.126 12694.756 4217.260 6447.805 11542.710
```

1 (d) (iii)

```
summary(x2)

##
## Iterations = 82081:86000
## Thinning interval = 1
## Number of chains = 4
## Sample size per chain = 4000
## 1. Empirical mean and standard deviation for each variable,
## plus standard error of the mean:
##
## Mean SD Naive SE Time-series SE
## mubeta1 12.190152 0.3632376 2.072e-03 2.929e-03
## mubeta2 0.028480 0.0082165 6.406e-05 7.396e-05
## signamubeta1sq 4.012323 1.1149435 0.028e-03 7.750e-02
## signamubeta2sq 0.001889 0.0005028 4.687e-06 7.959e-06
## sigmasy 0.056390 0.0042508 3.345e-05 3.949e-05
##
## 2. Quantiles for each variable:
##
## 2.5% 25% 50% 75% 97.5%
## mubeta1 1.147e+01 11.95215 12.192794 12.428588 12.993766
## mubeta2 1.239e-02 0.02301 0.02301 0.02301 0.044543
## signamubeta1sq 2.531e+00 3.29150 3.942758 4.468869 6.897768
## signamubeta2sq 9.538e-04 0.00139 0.001705 0.002114 0.003272
## sigmasy 4.870e-02 0.05345 0.050167 0.050113 0.065222
```

1 (d) (iv)

```
# reading the quantiles for mubeta2 from id111
exp(13*-0.01239)
## [1] 1.174767
exp(13*-0.04543)
## [1] 1.784359
```

An approximate 95% central posterior interval is $(1.174767, 1.784359)$ which is a much more restrictive interval than the first model.

1 (d) (v)

```
dic.samples(m2, 100000)
## Mean deviance: -16.63
## penalty BIC: 82
## Penalized deviance: 41.19
```

The effective number of parameters is 57.82 and the Plummer's DIC is 41.19.

1 (d) (vi)

The Plummer's DIC value for the first model was 42.97, which is higher than this model's DIC value. Therefore, this model would be preferred, although not by much.

I assumed that we can take X to just be a vector consisting of the i different years and the so the same X is used for each y.

